


Simulation Based Inference Library

Tutorial 0 – Getting Started

- Simulator – linear Gaussian - return the parameter + $1 + 0.1 * N(0,1)$
- Prior – box uniform, in all 3 parameters, $[-2,2]$
- Run 2000 simulations; $\theta = \text{prior.sample}()$, $x = \text{simulator}(\theta)$, $\text{inference} = \text{inference.append_simulations}(\theta, x)$, $\text{density_estimator} = \text{inference.train}()$, $\text{posterior} = \text{inference.build_posterior}(\text{density_estimator})$
- Then take the posterior, sample from it, simulate data from the posterior samples and plot
- We can assess the posterior for the known samples in this case
- Can also take the posterior, sample from it, simulate data from the posterior samples and assess the predictive power of the posterior
- Also track the log probability – compare the log prob of posterior sample with prior sample




Tutorial 1 – Amortized Posterior Inference on Gaussian Example

- Same basic example as before, but this time with amortization– ability to reevaluate the posterior for different observations without having to rerun inference
- Amortized posterior is one not focused on any one observation
- Two sets of observed data – run the inference on one of them to get a posterior, and then we can draw samples from the posterior given the second observation without rerunning inference
- These also match the ground truth parameters

Tutorial 2 – More Flexibility over Training Loop and Samplers

- 2D example from above, with prior now from $[-3,3]$
- Now we build our own density estimator – in this case we take an NSF of θ and x
- We use the Adam optimiser and optimise the loss from the density estimator to create our own custom training loop
- After this training, we can directly sample from the posterior given some observed data (need to be careful with batch dimensions)
- Can wrap this up into a direct posterior (like $\text{inference.build_posterior}$) to automatically reject samples outside prior bounds and compute Maximum-a-priori estimate
- We can also use custom data loaders in this way – can be helpful for more complex or larger datasets, such as images



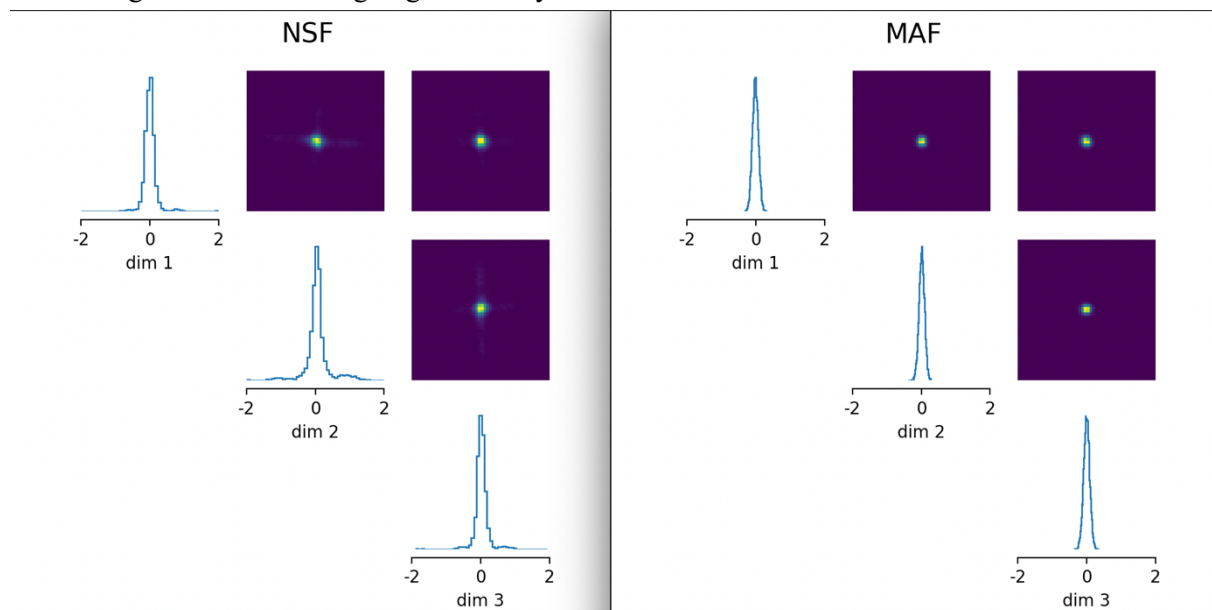
Simulation Based Inference Library

Tutorial 3 – Multi-round Inference

- So far – single round inference – draw parameters from prior, use simulator, train neural network to get parameter – can be inefficient if interested in only one observation
- Multi-round inference alleviates the problem – same as above, but after producing posterior it continues inference – this time it samples from the posterior conditioned on the observation, rather than the prior
- More efficient in the number of simulations, not amortized however, only accurate for this observation
- Beginning of the code – all the same; then create posteriors = [], start with proposal = prior and then update the proposal every time a new posterior is calculated, so each time you are training the density estimator with simulations from the most recent posterior

Tutorial 4 – Custom Density Estimators

- So far using standard density estimation methods – build_maf or build_nsf
- Can define custom density estimators – make a density estimator and then put it in the inference function; inference = NPE(...density_estimator=density_estimator)
- In this example – have done both NSF and MAF custom density estimators, compared the two of them using the same observed data in multi-round inference
- We can also change the hyperparameters of the density estimators in this way – for example, we can change number of hidden features, or number of transforms
- Can also build new density estimators completely from scratch – need to be careful with arguments when designing it this way

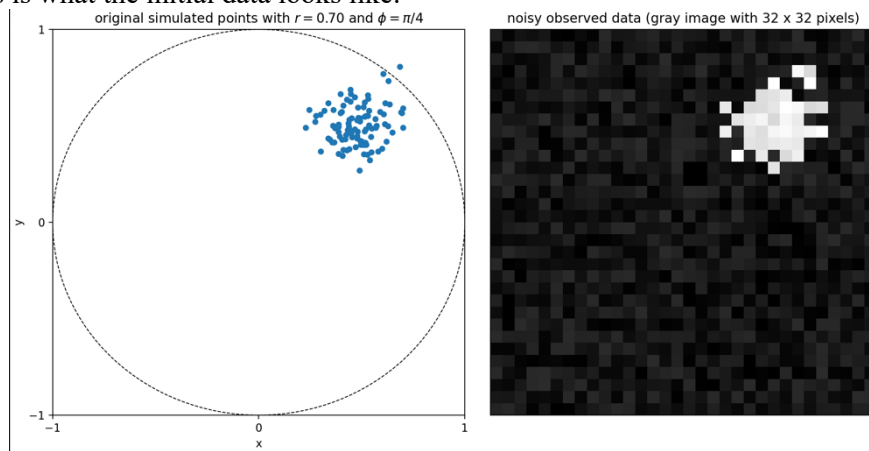


Tutorial 5 – Embedding nets for observations


- Summary statistics are very important – for high dimensional data especially. For complex data, this can be quite difficult to extract a priori, so you need to learn from the data which summary statistics to employ
- This is what embedding neural networks are for – the posterior can be written as $q_{\phi}(\theta|f_{\lambda}(x_o))$, where x_o are our observations, ϕ are the parameters of the conditional density estimator and λ the parameters of the embedding neural network
- Simulation goes to the embedding NN, then to the conditional DE; ϕ and λ are learnt jointly through the training
- In this example – our simulator takes (r, ϕ) in the 2D plane and produce Gaussian data centred at that point – then makes a greyscale image of the scattered points plus some noise as well.

Simulation Based Inference Library

- Then we take $32 \times 32 = 1024$ -dimensional data through a CNN embedding net to get compressed data, which we can then do our regular inference process on (this time with a MAF)
- This is what the initial data looks like:



- And then after inference we get the following results:






- Can also define custom embedding nets in the same way as custom density estimators – for this tutorial, defined custom CNN class called SummaryNet() and used that, can also import a more expressive CNN from the sbi library

Tutorial 6 – SBI with iid data

- If we have iid data, we might be interested in the posterior given all the data, $p(\theta | X = \{x_i\}_{i=1}^N)$
- When performing NPE, we cannot directly exploit iid assumption as we can do for NLE or NRE, so we take the full observation as an input, x . Thus, the embedding NN must be permutation invariant, and be able to handle a varying number of iid data points to be amortized, but when you have this, it is a fully amortized direct posterior
- To become invariant wrt number of trials, we need to run the simulator multiple times for individual parameter sets to get the training data – we get this done by interleaving in a for loop for this example
- To become permutation invariant, the neural net first learns embeddings for single trials and then performs a permutation invariant operation on those embeddings, e.g., by taking the sum or the mean

Simulation Based Inference Library

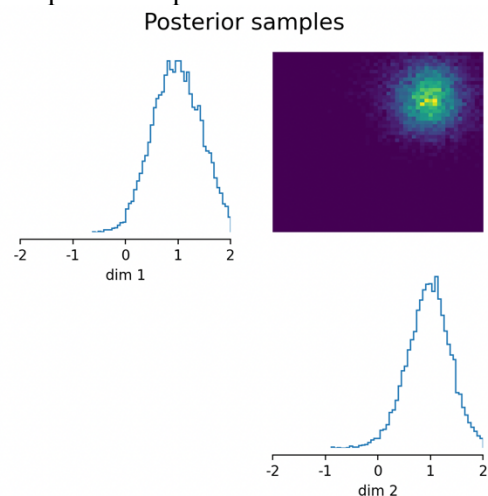
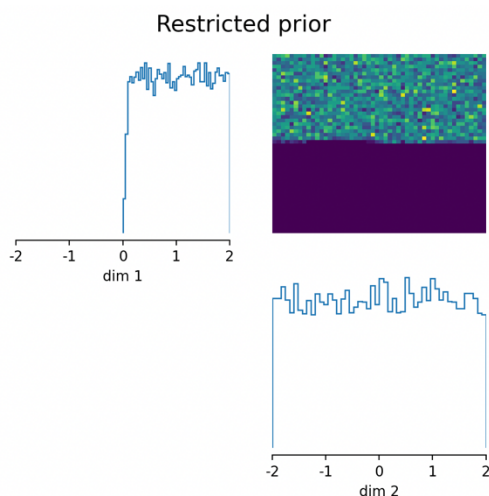
- We get lots of NaN values – this is from the missing trials – have to include this in training – and then we plot the posteriors learned from different numbers of iid data – we can see that the more data, the closer the posterior is centred around the true value



- We can easily obtain posteriors for many different observations, instantly, because NPE is fully amortize

Tutorial 7 – Efficient handling of invalid simulation outputs

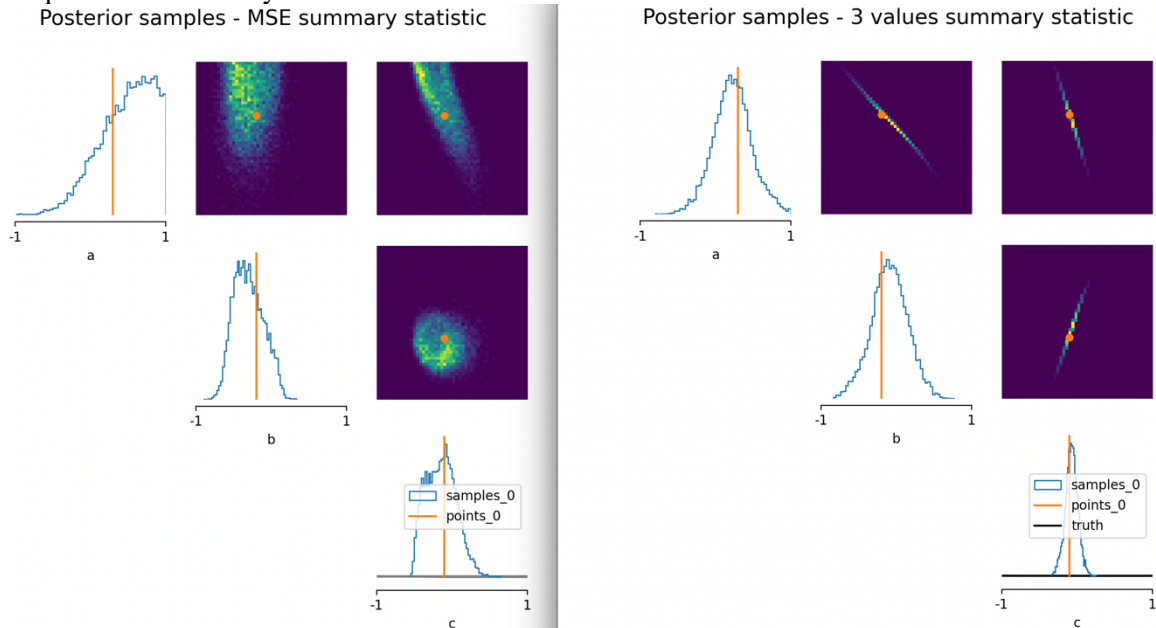
- Our simulators can give non-sensical outputs, these data should be rejected – however, this is inefficient – we do a lot of simulations and get little usable output
- We want the simulator to learn the parts of the parameter space which produce valid outputs – this can then discard any unsuitable parameters drawn from the start, making it more efficient
- Our simulator is defined in a way such that if the first parameter is less than 0 we return NaN, otherwise it's just Gaussian perturbation
- We set up the RestrictionEstimator – then use a classifier and train it so it can find the suitable parts of the parameter space to simulate from
- We then restrict our prior, and future posteriors, to this range, and then we simulate using this data – this time with all the data, rather than having to discard some
- Having trained the classifier to get a restricted prior, we then use it as before – all of a sudden, rather than getting many roughly 520 NaN results out of 1000 we get 0 (usually)
- Can see that the restricted estimator now has its first parameter prob = 0 below 0




Simulation Based Inference Library

Tutorial 8 – Crafting Summary Statistics

- Hand crafting summary statistics rather than learning them from the data itself as in tutorial 5
- Choosing 2 different summary stats for the data given and comparing their inference output – shows the importance of choosing the right summary stats
- The simulator produces some quadratic output, which is given by uniform prior
- Two summary stats considered – evaluate the function at 3 points, and the MSE between the observed and simulated data
- As observed – the 3-points summary statistic is more concentrated on true parameter, and the posterior is more symmetric about it – more useful information learned from this one



- Why? The summary statistics vary in terms of which simulations they learn from, as the diagram shows – the three points learn from all simulations that cover the observations, so the trained NN can interpolate the data – the MSE learns from simulations only to the right of the MSE so the NN has to extrapolate the data – first is much preferable




- Good strategy – create histograms like above, see how simulations cluster around observation – if they are off to one end of the simulations, not going to be as good

Tutorial 9 – Refining Posterior Estimates with Rejection Sampling

- SBI normally does sampling from prior and likelihood (via simulator) to get the posterior estimate – instead, we could use a mixture of a simulated estimate for the posterior, $q(\theta|x)$ and likelihood-based importance sampling to get an asymptotically exact estimate for the posterior, $p(\theta|x)$
- The idea is to think of $q(\theta|x)$ as a proposal, draw samples from the proposal $\theta_{y_0} \sim q(\theta|x)$ and then increase each sample with an importance weight $w_{y_0} = p(\theta|x) / q(\theta|x)$ – comes from Monte Carlo proof

Simulation Based Inference Library

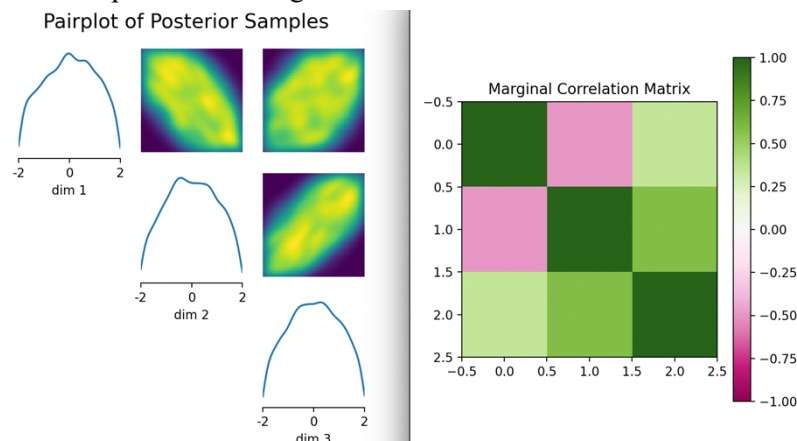
- Now we sample from $\theta_{y_o} \sim q(\theta|x)$ and the importance weights downweight samples where $q(\theta|x)$ oversamples $p(\theta|x)$ and vice versa
- We do this using the ImportanceSamplingPosterior function – this will do the importance sampling on the trained posterior, done in the same way as before
- We can compare the ground truth distribution, the NPE distribution and the NPE distribution with importance sampling – can see that the importance sampling is far more accurate



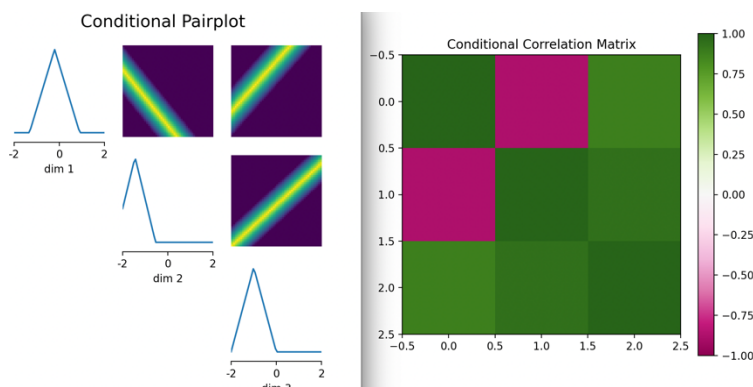
- Need some way of sampling from the tractable likelihood – not always helpful

Tutorial 10 – Variability and Compensation Mechanisms with Conditional Distributions

- Advantage of SBI – posterior captures all models that can reproduce experimental data
- Can check whether the parameters can be variable or have to be finely tuned, and also to find potential compensation mechanisms between model parameters – extract conditional distributions from inferred posterior
- Here is a trained posterior's marginals and the Pearson correlation coefficient matrix:



- As seen – little correlation (weak interactions), 1 and 2D marginals fill almost the entire parameter space – in 3D can look at the 3D plot, but inefficient for many dimensions, so we look at conditional marginals
- In the plots below, the diagonal marginals are all but 1 parameter constant, off diagonal is all but 2 parameters constant – can see the conditional correlation matrix as well



Simulation Based Inference Library